

Student Name: Iliadis Viktoras  
Student ID: 8180026

## Big Data Management Systems

### Assignment 2 - MapReduce/Hadoop

---

All the code and 2M generated data points can be found in the following [repository](#) that will host the code and all related files for all course exercises.

#### • Question 1 - Hadoop Installation

Hadoop was installed following the official documentation: <https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/SingleCluster.html>

```
hadoop version
```

```
Hadoop 3.3.5
```

```
Source code repository https://github.com/apache/hadoop.git -r  
706d88266abcee09ed78fbaa0ad5f74d818ab0e9
```

```
Compiled by stevel on 2023-03-15T15:56Z
```

```
Compiled with protoc 3.7.1
```

```
From source with checksum 6bbd9afcf4838a0eb12a5f189e9bd7
```

```
This command was run using
```

```
/home/mainuser/hadoop/hadoop-3.3.5/share/hadoop/common/hadoop-common-3.3.5.jar
```

Dependencies Installation

```
%pip install numpy matplotlib mrjob
```

#### • Question 2 - Data Generation

The following code generates a file with 2M datapoints (x,y) coordinate pairs. The generation is biased towards the creation of three clusters, by choosing a-priori three centers that you can find in the centroids.txt file. This is done by selecting a random distance r and a centroid randomly and then adding the generated distance r to the x and y coordinates of the centroid to create a new datapoint.

```
1 import random
2 import numpy as np
3
4 # Read the coordinates from the centroids.txt file
5 with open('centroids.txt', 'r') as f:
6     x1, y1 = map(float, f.readline().split())
7     x2, y2 = map(float, f.readline().split())
8     x3, y3 = map(float, f.readline().split())
9
10 # The three centers are stored in the variables x1, y1, x2, y2, x3, y3
11 # These values are read from the centroids.txt file, which contains the x and y
    coordinates
12 # of each of the three centers.
13
```

```

14 print("Initial centers:")
15 print(f"Center 1: ({x1}, {y1})")
16 print(f"Center 2: ({x2}, {y2})")
17 print(f"Center 3: ({x3}, {y3})")
18 print()
19
20 # Print the initial centers to the console for debugging purposes.
21
22 n = 2000000
23 data = []
24
25 # The code will generate 2 million (x, y) data points, which will be stored in the
    data list.
26
27 for i in range(n):
28     # Generate a random distance with some randomness added
29     r = abs(np.random.normal(loc=0, scale=3)) + abs(np.random.normal(loc=0, scale=2))
30
31     # Add the generated distance to one of the three centers chosen randomly
32     center = random.choice([(x1, y1), (x2, y2), (x3, y3)])
33     x, y = center[0] + r * np.random.normal(loc=0, scale=1), center[1] + r * np.
        random.normal(loc=0, scale=1)
34
35     # Add the generated (x, y) point to our data list
36     data.append((x, y))
37
38 # This loop generates 2 million (x, y) data points. For each point, a random distance
    r is generated using the
39 # normal distribution with mean 0 and standard deviation 3, with an additional normal
    distribution with mean 0 and
40 # standard deviation 2 added to introduce more randomness. Then, one of the three
    centers is chosen randomly,
41 # and the generated distance r is added to the x and y coordinates of the center. The
    resulting (x, y) point is
42 # added to the data list.
43
44 # Save the generated data to a text file
45 with open('clustered_data.txt', 'w') as file:
46     for x, y in data:
47         file.write(f"{x}, {y}\n")
48
49 # Finally, the generated data is saved to a text file named "clustered_data.txt",
    with each point on a new line.
50
51 print("Data generation complete")
52 print(n, "data points saved to 'clustered_data.txt'")
53
54 # Print a message indicating that the data generation is complete and how many data
    points were generated and saved.

```

```
%run data_generator.py
```

```

Initial centers:
Center 1: (10.0, 90.0)
Center 2: (1.0, 2.0)
Center 3: (78.0, 12.0)

```


```

Data generation complete
2000000 data points saved to 'clustered_data.txt'

```

Preview of the generated data points

```
1 # use matplotlib to plot the data
2 import matplotlib.pyplot as plt
3
4 # Read the data from the text file
5 data = []
6 with open('clustered_data.txt', 'r') as file:
7     for line in file:
8         x, y = line.strip().split(',')
9         data.append((float(x), float(y)))
10
11 # Plot the data
12 plt.scatter([x for x, y in data], [y for x, y in data])
13 plt.show()
```



- **Question 3 - Moving Data to HDFS**

We place the datapoints in the hdfs using the following command

```
hdfs dfs -put clustered_data.txt clustered_data.txt
```

And we also run

```
hdfs dfs -rm -r clustered_data.txt
```

beforehand to remove any old file if it exists.

- **Question 4 - Execute 1 iteration of K-means**

Utilizing Hadoop Streaming and the code in kmeansiteration.py, one iteration of K-means is executed. In the end, the centroids are printed to the console.

```
1 from mrjob.job import MRJob
2 from mrjob.protocol import RawValueProtocol
3 import math
4
5 class KMeansIteration(MRJob):
6     # Method to configure command line arguments
7     def configure_args(self):
8         super(KMeansIteration, self).configure_args()
9         self.add_passthru_arg('--centers', help='comma-separated list of centers')
10
11     # Method to initialize the mapper and load the center values from the command
12     # line arguments
13     def mapper_init(self):
14         centers = list(map(float, self.options.centers.split(',')))
15         self.centers = [(centers[i], centers[i+1]) for i in range(0, len(centers), 2)]
16
17     # Mapper method to calculate the closest center for each data point and emit the
18     # key-value pair
19     def mapper(self, _, line):
20         data = list(map(float, line.split(',')))
21         distances = [math.sqrt((data[0]-center[0])**2 + (data[1]-center[1])**2) for
22                     center in self.centers]
23         closest_center_index = distances.index(min(distances))
24         yield closest_center_index, data
25
26     # Reducer method to calculate the new center for each cluster and emit the key-
27     # value pair
28     def reducer(self, key, values):
29         new_center = [0.0, 0.0]
30         count = 0
31         for value in values:
32             new_center[0] += value[0]
33             new_center[1] += value[1]
34             count += 1
35         new_center[0] /= count
36         new_center[1] /= count
37         yield key, new_center
38
39     # Set the input protocol to RawValueProtocol
40     INPUT_PROTOCOL = RawValueProtocol
41
42 if __name__ == '__main__':
43     KMeansIteration.run()
```

```
python kmeans_iteration.py -r hadoop --hadoop-streaming-jar
/home/mainuser/hadoop/hadoop-3.3.5/share/hadoop
/tools/lib/hadoop-streaming-3.3.5.jar
hdfs://localhost:9000/user/mainuser/clustered_data.txt
--centers 10,10,50,70,90,30
```

No configs specified for hadoop runner

Looking for hadoop binary in /home/mainuser/hadoop/hadoop-3.3.5

```

/bin...
Found hadoop binary: /home/mainuser/hadoop/hadoop-3.3.5
/bin/hadoop
Using Hadoop version 3.3.5
Creating temp directory /tmp/kmeans_iteration.mainuser.
20230422.160133.633377
uploading working dir files to hdfs:///user/mainuser/tmp/mrjob
/kmeans_iteration.mainuser.
20230422.160133.633377/files/wd...
Copying other local files to hdfs:///user/mainuser/tmp/mrjob
/kmeans_iteration.mainuser.
20230422.160133.633377/files/
Running step 1 of 1...
packageJobJar: [/tmp/hadoop-unjar7949111280189398112/] []
/tmp/streamjob2377951395818703915.jar tmpDir=null
Connecting to ResourceManager at /0.0.0.0:8032
Connecting to ResourceManager at /0.0.0.0:8032
Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/mainuser
/.staging/job_1682172351424_0025
Total input files to process : 1
number of splits:2
Submitting tokens for job: job_1682172351424_0025
Executing with tokens: []
resource-types.xml not found
Unable to find 'resource-types.xml'.
Submitted application application_1682172351424_0025
The url to track the job: http://LAPTOP-E3S2CMV1.:8088
/proxy/application_1682172351424_0025/
Running job: job_1682172351424_0025
Job job_1682172351424_0025 running in uber mode : false
  map 0% reduce 0%
  map 33% reduce 0%
  map 53% reduce 0%
  map 76% reduce 0%
  map 100% reduce 0%
  map 100% reduce 89%
  map 100% reduce 100%
Job job_1682172351424_0025 completed successfully
Output directory: hdfs:///user/mainuser/tmp/mrjob
/kmeans_iteration.mainuser.20230422.160133.633377/output
Counters: 54
File Input Format Counters
Bytes Read=75664139
File Output Format Counters
Bytes Written=126
File System Counters
FILE: Number of bytes read=87660049
FILE: Number of bytes written=176162026
FILE: Number of large read operations=0
FILE: Number of read operations=0

```

FILE: Number of write operations=0  
HDFS: Number of bytes read=75664351  
HDFS: Number of bytes read erasure-coded=0  
HDFS: Number of bytes written=126  
HDFS: Number of large read operations=0  
HDFS: Number of read operations=11  
HDFS: Number of write operations=2  
Job Counters  
Data-local map tasks=2  
Launched map tasks=2  
Launched reduce tasks=1  
Total megabyte-milliseconds taken by all map tasks=56822784  
Total megabyte-milliseconds taken by all reduce tasks=18595840  
Total time spent by all map tasks (ms)=55491  
Total time spent by all maps in occupied slots (ms)=55491  
Total time spent by all reduce tasks (ms)=18160  
Total time spent by all reduces in occupied slots (ms)=18160  
Total vcore-milliseconds taken by all map tasks=55491  
Total vcore-milliseconds taken by all reduce tasks=18160  
Map-Reduce Framework  
CPU time spent (ms)=63050  
Combine input records=0  
Combine output records=0  
Failed Shuffles=0  
GC time elapsed (ms)=2406  
Input split bytes=212  
Map input records=2000000  
Map output bytes=83660043  
Map output materialized bytes=87660055  
Map output records=2000000  
Merged Map outputs=2  
Peak Map Physical memory (bytes)=506572800  
Peak Map Virtual memory (bytes)=2598981632  
Peak Reduce Physical memory (bytes)=282497024  
Peak Reduce Virtual memory (bytes)=2602565632  
Physical memory (bytes) snapshot=1221701632  
Reduce input groups=3  
Reduce input records=2000000  
Reduce output records=3  
Reduce shuffle bytes=87660055  
Shuffled Maps =2  
Spilled Records=4000000  
Total committed heap usage (bytes)=1286078464  
Virtual memory (bytes) snapshot=7697793024  
Shuffle Errors  
BAD\_ID=0  
CONNECTION=0  
IO\_ERROR=0  
WRONG\_LENGTH=0  
WRONG\_MAP=0

```

WRONG_REDUCE=0
job output is in hdfs:///user/mainuser/tmp/mrjob
/kmeans_iteration.mainuser.20230422.160133.633377/output
Streaming final output from hdfs:///user/mainuser/tmp/mrjob
/kmeans_iteration.mainuser.20230422.160133.633377/output...
0 [1.0127311415976215, 2.0063022791532448]
1 [9.998596586236912, 89.99545027492118]
2 [78.01356094767394, 12.009210995449154]
Removing HDFS temp directory hdfs:///user/mainuser/tmp/mrjob
/kmeans_iteration.mainuser.20230422.160133.633377...
Removing temp directory /tmp
/kmeans_iteration.mainuser.20230422.160133.633377...

```

### • Question 5 - Complete k-means

Centroids in centroids.txt were manually changed before the execution to demonstrate more iterations.

Utilizing the code in kmeans.py, the complete K-means algorithm is executed. The centroids are printed to the console after each iteration, as well as the Average Distance from the centroids of the previous iteration.

```

1 from kmeans_iteration import KMeansIteration
2 import math
3
4
5 # Function to run the KMeans algorithm on the input data
6 def run_kmeans(data_path, centers, threshold=0.01, max_iterations=10):
7     # Store the initial centers to check for convergence later
8     old_centers = centers
9
10    # Loop through the maximum number of iterations
11    for i in range(max_iterations):
12        print(f'Iteration {i + 1}')
13
14        # Set up the arguments for the KMeansIteration MRJob
15        job_args = [
16            data_path,
17            '—centers', ' ', ' '.join(map(str, centers)),
18            '—r', 'hadoop',
19            '—hadoop-streaming-jar',
20            '/home/mainuser/hadoop/hadoop-3.3.5/share/hadoop/tools/lib/hadoop-
streaming-3.3.5.jar'
21        ]
22
23        # Create a new instance of the KMeansIteration MRJob and run it
24        job = KMeansIteration(args=job_args)
25        with job.make_runner() as runner:
26            runner.run()
27
28        # Collect the new centers from the output of the MRJob
29        centers = []
30        for _, center in job.parse_output(runner.cat_output()):
31            centers.extend(center)
32            print(center)
33
34        # Check for convergence
35        if converged(old_centers, centers, threshold):
36            print('Converged')

```

```

37         break
38
39         # Update old_centers to be the current centers for the next iteration
40         old_centers = centers
41
42
43 # Function to check for convergence between the old and new centers
44 def converged(old_centers, new_centers, threshold):
45     total_distance = 0
46     for i in range(0, len(old_centers), 2):
47         # Calculate the Euclidean distance between the old and new centers
48         total_distance += math.sqrt(
49             (old_centers[i] - new_centers[i]) ** 2 + (old_centers[i + 1] -
50             new_centers[i + 1]) ** 2)
51         average_distance = total_distance / (len(old_centers) / 2)
52         print(f'Average distance: {average_distance}')
53         print(f'Threshold: {threshold}')
54         return average_distance < threshold
55
56 if __name__ == '__main__':
57     data_path = 'hdfs://localhost:9000/user/mainuser/clustered_data.txt'
58
59     # Read the coordinates from the centroids.txt file
60     with open('centroids.txt', 'r') as f:
61         x1, y1 = map(float, f.readline().split())
62         x2, y2 = map(float, f.readline().split())
63         x3, y3 = map(float, f.readline().split())
64
65     initial_centers = [x1, y1, x2, y2, x3, y3]
66
67     # Run the KMeans algorithm
68     run_kmeans(data_path, initial_centers)

```

python kmeans.py

Iteration 1

No configs specified for hadoop runner  
[76.59013482611823, 12.074173840865583]  
[1.126995594698583, 1.971656929112909]  
[16.461721650478808, 83.91995985391885]  
Average distance: 34.435277361403045  
Threshold: 0.01

Iteration 2

No configs specified for hadoop runner  
[78.00541564910183, 12.007950356032087]  
[0.9999430159528663, 2.0018522657735165]  
[9.996448832210588, 89.99503318486315]  
Average distance: 3.473025362811198  
Threshold: 0.01

Iteration 3

No configs specified for hadoop runner  
[78.0054829362083, 12.008062334929722]  
[1.0000850823919099, 2.0019539984736503]  
[9.996395522177313, 89.99523047792616]  
Average distance: 0.00016991463780164826



Threshold: 0.01  
Converged